Odd-prime kernels: complete measurement inputs

Human Compression project

Expanded Topic F; input proposal only

This boundary compares two complete originals with one complete common paper and two complete applications. It preserves 25 L014 and 38 L007 scientific source dispositions. It computes input lengths, not compression factors.

L014: Hou, Feng, Leemans and Qu, *String C-groups of order* $4p^m$, arXiv:2407.10388v1, 15 July 2024, twelve pages. **L007:** Kong, Feng, Hou, Leemans and Qu, *Chiral 3-polytopes of order* $2p^m$, accepted 19 January 2026, DOI 10.26493/1855-3974.3716.b64, twenty-two pages. Every physical source page is included, with its title, abstract, formulas, tables, bibliography and contact information.

Complete document	Pages	Raw NFC characters
Original L014	12	25,574
Original L007	22	44,251
Common paper U	7	15,876
L014 application A1	5	11,504
L007 application A2	6	15,894
Complete program appendix	9	22,466

Characters here mean non-whitespace characters in the untrimmed extraction. This primary coordinate was fixed before extraction. The six PDFs occupy 61 pages. No old common paper, application or program is an unpaid premise.

A later written comparison must pay U+A1+A2 in full: eighteen pages, including reference-only endings. A separate program-inclusive comparison adds the whole nine-page appendix. The alternative raw representation instead adds all nine complete program files, 25,501 bytes. These representations must not be added together. Each conditional source/application comparison requires prior acquisition of the entire U and all nine declared programs; its remaining local cost is all of A1 or all of A2.

Programme roles. All nine files are finite validation, not premises of uniform theorems. They include the three inherited L014 checks, two current carry and regular-family checks, two exploratory scans, and two inherited central-lift checks. Complete files are retained even when they are diagnostic. The appendix's nine full `VerbatimInput` directives and a text-inclusion check account for every file once. Continuation marks are removed only in that printing check; every printed mark remains in every measured coordinate.

Literal extraction and visible glyphs

All five established units are recorded under NFC and NFKC for untrimmed text, selected text and calibrated selected text: 180 PDF cells. Source L014's arXiv banner and L007's 22 accepted-manuscript watermarks remain in the raw primary. The selected variants remove exact page furniture only: repeated running headers, those watermarks and isolated replacement page numbers. They preserve source content, original hyphens and extraction spacing.

The original PDF byte events resolve to 1,140 used classes across 50 actual font dictionaries. Dictionary identity, effective encoding, font stream and outline evidence are bound separately. Sixty-three mathematical and icon outlines were freshly inspected. Thirty-seven classes receive 780 occurrence-bound sensitivity repairs, covering mathematical alphabets, parentheses, brackets, bracket pieces, sums, products, tildes and one mapsto component class. The latter uses the declared arrow proxy for that component; the adjoining arrow is still counted.

Two licence icons and four horizontal-brace component classes in L007 stay literal encoded proxies. Their visual structure is not reconstructed as text. The source logos, table rules, formula placement and full graphical complexity remain unknown, not zero. Ordinary variable font styling is not re-encoded. The four replacement PDFs retain their Unicode Type0 text unchanged; no Type1 source-outline claim is made for those fonts.

The PDF visitor emits one duplicate, fontless watermark callback on every L007 page. The narrowly bound helper ignores that duplicate only for metadata alignment. The returned extraction is byte-for-byte the unchanged standard `pypdf` text, and the actual PDF stream still contains each watermark once. No mathematical source character is dropped by this repair.

Complete scope and separate costs

Four complete operational drivers total 33,222 bytes, and three complete Lean files total 9,398 bytes. Together with the nine programs these yield 160 raw metric cells. The twelve scoped declarations comprise nine reused and three new generic statements. No particular source construction, complete classification or geometric correspondence is formally proved by these files. These operational and formal costs are separate from paper/program comparisons.

The corrected scientific boundary preserves the full L014 contract and all L007 parameter ranges. In L007's Remark 1.3 the fourth regular family exists exactly for $p = 3$; the Section 6 quotient argument is repaired. These corrections do not rewrite the measured primary. The previously written cyclic lift is charged in U and has a distinct fresh root adjudication. No new question or historical-priority claim is made.

Complete external proofs of the retained Cunningham–Pellicer, Huppert, Schmidt/Iwasawa and classical inputs remain unpriced. Scientific acceptance is reused, with all 61 page views bound; six source formula/icon pages were re-viewed for calibration. The measurement author was the scientific reviewer, with the inherited lift role disclosed. A distinct input audit is still required. No mathematical program or Lean theorem is rerun by this measurement packet.