Inputs for two locally projective quotient papers

This proposal fixes the complete input boundary. It does not calculate a compression factor. The mathematical migration has passed distinct root review; independent acceptance of these extraction and calibration choices is still required. All 52 inherited source records and their qualifications are retained.

Complete papers

L084 is the two-page publisher erratum of Hartley and Leemans (2006). L095 is their fifteen-page arXiv v1, dated 5 July 2003. The unseen journal edition is not substituted. Titles, abstracts, addresses, references, the version banner, and every historical table remain included. Calibration does not repair the source's erroneous historical counts.

The whole common acquisition comprises the seven-page proof and five-page census companion. It contains the marked universal-group proof needed by both complete applications. No application obtains a missing premise free from its sibling. The previous fourteen-page written arrangement remains immutable provenance; it is not an extra current cost.

Input	Pages	Raw	Selected	Calibrated
L084: publisher erratum (2006)	2	2,095	2,071	2,071
L095: arXiv v1 (2003)	15	24,440	24,419	24,419
Whole common proof	7	19,540	19,533	19,533
Whole common census tables	5	5,970	5,965	5,965
Complete L084 application	2	4,859	4,857	4,857
Complete L095 application	3	7,112	7,109	7,109
All six programs, printed once	6	22,701	22,695	22,695

The three integer columns count NFC characters excluding whitespace. The selected variants remove only the erratum's second-page running header and isolated trailing page numbers. The first-page publication metadata is retained. In total, seven PDFs contain forty physical pages. Physical extent and extractable text are distinct coordinates.

A later written comparison uses both originals against the entire common bundle plus both applications. A per-original conditional comparison uses that original and its complete application *after acquiring the whole common proof and tables, all six required programs, all required comparison data, and retained external inputs*. No cross-topic discount is proposed.

Complete programs and data

All six programs are required finite calculations. They include the marked group, presentation bound, local geometry, full maximal-subgroup and quotient censuses, and exact reconciliation. The shared marking file is charged once. The six-page appendix prints all six files without line pruning.

Whole raw category	Bytes	NFC nonspace
Six whole required programs	24,255	22,104
Whole historical comparison data	39,150	28,574
Three complete operational drivers	32,220	29,370

The complete historical comparison file contains all 145 literal rows. It is required by reconciliation, although it is not a subgroup-enumeration premise. Either the whole printed appendix *plus this raw data file*, or the six whole raw programs *plus the same data file*, may provide a later programme comparison. Printed and raw copies of the same code are alternatives; they are never added together.

All eleven output/evidence files are separately bound, totalling 3,470,182 bytes. This includes the complete computed subgroup representatives. They are regenerated outputs, not an unseen catalogue input. The three full operational drivers are reported separately. Thirteen selected installed algorithm/model pins are preserved, without claiming they measure the complete software installation. No Lean module or formal declaration is part of this science bundle.

Actual glyphs and sensitivity

The audit inspected all seventeen source pages and all 502 actual source glyph classes. Eight dictionary-and-byte classes require 61 calibrated occurrences: mathematical calligraphic capitals, one correspondence envelope, two proof-end squares and four mapsto tails. Each repair binds the actual font dictionary, effective encoding, encoded byte, event and page context. The mapsto tail remains a separate event from its adjacent arrow. No source formula is emended, no two events are merged, and no graphical proxy is used. Both raw variants retain the literal extracted characters.

The parser explicitly handles three input-specific cases: a name-keyed CFF font's omitted charset operator has its ISOAdobe default; an explicit MacRoman encoding is respected before PDF differences; and the standard unembedded Times-Roman banner uses its default encoding. Fonts sharing a name are never assumed to share an encoding. The current five replacement PDFs use unchanged Unicode extraction, with all twenty-three pages visually reviewed in the science migration.

L095 contains inline image rules, principally table borders; L084 has publisher rules. Their presence is inventoried separately. No absent mathematical diagram was identified, but table layout and graphic construction are not measured by text or by an image count. Full retained classical proof, runtime, historical software, graphics/layout and full formalization costs remain unknown, not zero. The unavailable historical `g.txt` is not a premise of the replacement computation.

The data retain all five established units (whitespace tokens, alphanumeric units, characters without whitespace, characters with collapsed whitespace, and UTF-8 bytes with collapsed whitespace), both NFC and NFKC, and all three extraction variants. There are 210 PDF metric cells and 100 raw-file metric cells. Printed hyphens and ligatures are not silently modernized. Exact text hashes, raw sensitivity and every selected cut are preserved.

The portable recount requires the pinned original PDFs and accepted science files. It runs no GAP calculation. Separate controls test omitted full-paper and data costs, body truncation, font/byte substitution, and the actual parser defaults. No primary text, font program, outline image or page raster is redistributed in this input packet.